

Lumielle

Academic Assistant

Academic Paper Writing Assistant

USER GUIDE

macOS · Windows

Open-source · MIT License



CYBER ANGEL

CYBER ANGEL
光 ← 未来
-XIA
X-01

—X
XIX
+

CODE: X-01

01

What Lumielle does

Lumielle keeps your research materials, references, outline, chapter drafts, and Word export in one writing workflow. You can return to an earlier chapter at any time; once its chapter memory is updated, later generation uses the revised information.

Common situation	How Lumielle handles it
Long chats lose track of earlier chapters	Each written chapter has a compact memory used when other chapters are generated
An earlier chapter changes but later writing still follows the old version	Edit it and click “Save & update memory”; later generation uses the new memory
Important findings appear late in a long paper	The complete extractable document is processed before its profile and evidence are built
Task briefs, methods, data, and background are mixed together	The Research Foundation stores them by role and controls which material may support chapter writing
References drift away from the chapters that use them	Each chapter can bind specific literature; chapter writing selects evidence from those bound sources
The draft still needs manual Word formatting	Choose a built-in or shared formatting preset and export to Word

Generated text still needs review

Lumielle reduces common risks such as forgetting context, repetition, scope drift, and citation mismatch, but the author remains responsible for data, facts, citations, and the final manuscript.

1	2	3	4	5	6	7
Topic	Research material	Literature	Outline	Chapters	Review	Word export

A practical order is to prepare the topic and key materials first, build the outline, check chapter structure and literature bindings, then start drafting. On a first run, generate one chapter before committing to a full draft.

02

Install and start

Windows

1. Locate the downloaded ZIP archive.
2. Right-click the ZIP and choose “Extract All”. Do not run the application from inside the ZIP preview.
3. Open the extracted folder and double-click “Launch.bat”.
4. The first launch may take a few minutes. Keep the launcher window open while setup is running.
5. When the Lumielle page opens in your browser, the application is ready. Continue to use the same extracted folder in future sessions.

macOS

1. Double-click the ZIP file to extract it.
2. Open the extracted folder and double-click “Launch.command”.
3. If macOS blocks the first launch, Control-click / right-click the launcher and choose “Open”. If needed, use System Settings → Privacy & Security → Open Anyway for that launch.
4. The first launch may take a few minutes. Continue after the browser page opens.

Keep the working data folder

After you start using Lumielle, the data folder stores configuration, drafts, literature, and project data. Do not delete or overwrite it unless you know exactly what you are removing.

Connect an AI model

Lumielle manages the paper structure, research material, literature, chapter memory, and export workflow. Text generation and document analysis still require a large language model. An API is simply the connection Lumielle uses to call that model directly.

Lumielle uses OpenAI-compatible endpoints. DeepSeek is a practical choice for Chinese academic writing, but other compatible providers can also be used with their own Base URL, API Key, and model.

DeepSeek App/Web and the API are separate services

Even if you use the DeepSeek app or web chat, Open Platform API access depends on the API account status and balance. API usage is billed separately by the provider.

03

Create a DeepSeek API Key

DeepSeek Open Platform: <https://platform.deepseek.com>

API Keys page: https://platform.deepseek.com/api_keys

The steps below follow the current DeepSeek Open Platform. Button wording may change slightly over time; look for the API Keys section if the page layout changes.

1. Open the DeepSeek Open Platform and sign in. An existing DeepSeek account can normally be used for the Open Platform as well.
2. Open “API Keys” from the left-side navigation, or use the direct API Keys link above.
3. Choose “Create new API key”. If a name is requested, use “Lumielle” or any label you will recognize later.
4. After creation, the platform shows a string beginning with “sk-”. The entire string is the API Key; do not copy only the first or last few characters.
5. The complete key is shown only when it is created. Copy and store it immediately. If it is lost, disable/delete the old key and create a new one.
6. Check your Open Platform balance. If there is not enough balance, Lumielle may report “insufficient balance” during the connection test.

API Key example

A key typically looks like sk-xxxxxxxxxxxxxxxxxxxx. This is only an example. Never post your real key in public screenshots, chats, forums, or shared documents.

DeepSeek’s current OpenAI-compatible Base URL is <https://api.deepseek.com>. The model list fetched inside Lumielle is the authoritative list for your account. For Chinese writing, deepseek-flash is a good starting point; deepseek-v4-pro can be tried when it is available and you want the stronger model.

04

Finish model setup in Lumielle

1. Open “LLM Configuration & Multi-Model Cross Discussion” and select the LLM Configuration Matrix.
2. Expand “Add New LLM Configuration”. Give the profile a recognizable name such as “DeepSeek Writing”.
3. Set Base URL to <https://api.deepseek.com>. New installations already use this as the default.
4. Paste the complete sk-... value into API Key. The field hides the key while you type; that is expected.
5. Click “Fetch Model List”. A model dropdown appears after the provider returns the models available to your account.
6. Select a model. deepseek-flash is a good first choice for Chinese writing; deepseek-v4-pro can be selected when available.
7. Click “Test Model Connection”. Lumielle sends one very small real chat-completion request to confirm the Base URL, key, and model can actually be called.
8. When the green “Model connection is working / Successfully called” message appears, click “Save New Configuration”.
9. Make sure the profile is marked “Currently Active”. If not, click “Activate”.

What the green result means

The selected key, model, and endpoint completed a real request at the time of testing. Later changes to balance, network, provider status, or model availability can still require another test.

Message	Typical cause	What to do
Authentication failed	The key is incomplete, wrong, or no longer valid	Copy the correct full key again; create a new key if the old one was lost
Insufficient balance	The API account has no usable credit	Check the provider balance / top-up page
Model unavailable	The saved model is gone or not enabled for the account	Fetch the model list again and select from the returned list
Rate limited	The provider is limiting request frequency	Wait and try again
Timeout / connection error	Network or Base URL issue	Check the network and Base URL, then retry

Other AI providers

Any provider with an OpenAI-compatible endpoint can be configured in the same general way. Use the Base URL, API Key, and model information supplied by that provider.

05

Topic and Research Foundation

Before writing, set the research topic and estimated total word count in the sidebar. The topic is used later for literature relevance, outline planning, and chapter generation.

What belongs in the Research Foundation

Use the Research Foundation for information that belongs to your own project: task briefs, background, objectives, definitions, methods, experimental or survey data, findings, and limitations.

Section	Typical content	Eligible for chapter writing by default
Research Background	Context, motivation, problem background	Off
Assignment / Requirements	Supervisor task brief, project requirements	Off
Objectives & Questions	Aims, research questions, hypotheses	On
Theory / Concepts / Variables	Definitions, variables, theoretical framework	On

Section	Typical content	Eligible for chapter writing by default
Research / Experimental Methods	Experiments, surveys, interviews, models, analysis methods	On
Experimental / Survey Data	Raw data, tables, statistics	On
Results / Findings	Observed results and core findings	On
Constraints / Boundaries	Sample limits, scope, project constraints	On

“Allow in chapter writing” is not a quality slider

When enabled, relevant facts from that section may be selected for the chapters that need them. It does not mean the whole section is inserted into every chapter. Disabled sections can still influence planning and outline generation.

Add material

- Each section accepts PDF, DOCX, and TXT files.
- You can also paste text directly.
- If you are unsure where something belongs, use the unclassified / smart filing area, review the suggested section, and confirm it before saving.
- Long sources are processed from their complete extractable text rather than only the beginning.

Scanned PDFs

If a PDF contains only images and no selectable text, v1.1.0 does not include OCR. Convert the scan to a text-extractable PDF or DOCX before importing it.

06

Import and organize literature

The literature library stores papers and sources you want to use in the manuscript. PDF, DOCX, TXT, and pasted text are supported.

After import, Lumielle processes the complete extractable text and builds a profile containing the research question, methods, sample or data, major findings, limitations, and summary. It also stores source-linked evidence that can later support chapter writing.

How to read the rating

The star rating mainly reflects relevance and usefulness to the current research topic. It is not a simple score of whether the paper is “good” or “bad”. Methodological quality, evidence strength, and limitations are assessed

separately.

Chapter bindings

Outline generation can bind literature to different chapters, and you can edit those bindings manually. When a chapter is generated, literature evidence is selected only from sources already bound to that chapter instead of from the entire library.

About fabricated references

Outline citations and chapter bindings are selected from literature records that already exist in your library. This reduces the risk of inventing a new reference that is not in the library. Facts, quotations, page references, and the final bibliography still require human verification.

Clearing the library

Clearing literature records and deleting the imported source files are separate options. You can remove the library records while keeping the source files, or explicitly choose to remove both.

07

Generate and review the outline

Build the logic outline after the Research Foundation and core literature are ready. The outline controls chapter titles, order, target length, experiment flags, and literature bindings.

1. Confirm the topic and total target length.
2. Check that important methods, data, and findings are saved in the Research Foundation.
3. Import the main literature you expect to use.
4. Generate the logic outline.
5. Review chapter titles, descriptions, target word counts, and literature bindings.
6. Edit the outline where necessary before drafting chapters.

Do not rush into full batch generation

On a first project, generate one chapter and inspect the style, references, and use of your materials. Correcting the workflow early is usually easier than revising an entire first draft.

Existing Word manuscripts

If you already have a DOCX manuscript, Manuscript Reverse Engineering can read the complete document structure and build a logic tree. If reverse engineering fails, the previous outline is kept.

08

Chapter writing and the Context Sandwich

Lumielle does not treat the paper as one endlessly growing chat. When a chapter is generated, the request combines the paper outline, memories of written chapters, the current chapter requirements, downstream boundaries, and relevant research facts and literature evidence.

Used for the current chapter	Purpose
Global paper outline	Places the chapter within the whole paper and keeps its scope clear
Memories of earlier written chapters	Carries forward definitions, data, conclusions, and prior argumentation
Current chapter requirements	Provides the title, scope, description, and target length
Downstream chapter boundaries	Reduces writing content that belongs in later chapters
Relevant research facts and literature evidence	Provides project-specific and source-linked support for the chapter

This chapter context is called the “Context Sandwich”. It reduces common long-form problems such as forgetting earlier material, repetition, scope drift, and inconsistency after revising a middle chapter, but it cannot guarantee that every conflict is automatically eliminated.

After you edit a chapter that was already written

Click “Save & update memory”. Editing the text alone does not automatically rebuild that chapter memory. Other chapters that were already generated are not rewritten automatically; if the revision affects later chapters, return to those chapters and review or regenerate them.

1	2	3	4
Return to chapter	Edit text	Save & update memory	Continue writing

Single chapter vs batch generation

- Single chapter: better for close control and review; recommended for a first project.
- Batch generation: useful for producing a complete first draft quickly; every chapter can still be revised afterward.
- In either mode, saved chapter text remains editable. Rebuild chapter memory when a revision should affect later generation.

09

Charts, images, and AIGC tools

Charts

Charts are rendered from validated structured chart specifications instead of executing temporary Python generated by the model. Real data should be used when available; illustrative values are only allowed when you explicitly authorize illustrative data.

Images

You can attach your own images to chapters. The Word export preserves linked images and charts. Check resolution, captions, and source attribution before submission.

AIGC detection and rewriting

AIGC detection is a writing-style aid, not a statistically guaranteed forensic result. If detection is unavailable, Lumielle reports it as unavailable rather than showing a misleading 0%. Rewriting can reduce mechanical phrasing, but you must still check that technical meaning has not changed.

10

Word export and formatting presets

After the manuscript is complete, choose a formatting preset in the export area and generate a Word document. Formatting rules are validated and applied by the local renderer; you do not need to write code.

Option	When to use it
General Academic	Common academic layout: A4, academic margins, three-line tables, page numbers, and captions
Legacy Compatible	A layout closer to the export behavior of earlier Lumielle versions
School / department JSON preset	Someone has already created and shared the required institutional formatting rules
Create a preset from written requirements	You only have instructions such as "12 pt body, 1.5 line spacing..."
DOCX template + formatting preset	You need to keep existing template content/styles and apply supported formatting rules

You do not need to write JSON

JSON is only the file format used to store a formatting preset. A shared “University-Thesis.json” preset can be imported directly. When you generate a preset from natural-language requirements, review the preview, save it, then select it for export.

v1.1.0 supports A4 page size, margins, body and three heading levels, paragraph spacing, page-number start, grid / three-line tables, figure and table captions, image width controls, heading pagination behavior, citation remapping, and template-preserving export. Institution-specific covers, declarations, complex headers, or special bibliography styles may still require additional adjustment.

Always review the exported Word file

Pagination can change with fonts, Office versions, and templates. Before submission, check page numbers, the table of contents, figure/table page breaks, equations, references, and the institution’s current requirements.

11

Recommended workflow

1. Start Lumielle and confirm the model connection test passes.
2. Set the research topic and total target length.
3. Add your task brief, methods, data, and findings to the Research Foundation.
4. Import the core literature and let analysis finish.
5. Generate the logic outline and review chapters, target lengths, and literature bindings.
6. Generate 1-2 chapters first and check writing style and citations.
7. Continue with single-chapter or batch generation.
8. When an earlier chapter changes, return to it, edit, and click “Save & update memory”.
9. Review structure, data, references, charts, and AIGC-style warnings across the manuscript.
10. Choose a formatting preset and export to Word.
11. Perform final human proofreading before submission.

Existing chapters are not silently rewritten

Updating chapter memory changes the context used for later generation, but it does not secretly overwrite chapters that are already saved. If a Chapter 2 revision affects a Chapter 5 conclusion, return to Chapter 5 and review it.

12

Frequently asked questions

Do I have to upload a task brief?

No. A task brief is only one type of research material. You can work from a topic plus your existing methods, data, background, or other project information.

Should every Research Foundation section be allowed in chapter writing?

No. Methods, data, and findings are often useful in chapters; administrative instructions or background material may be better left off. Disabled sections can still inform planning.

Will a long paper be analyzed only from the beginning?

No for text-extractable documents in v1.1.0. The complete extractable text is processed in order before the profile and evidence are built. Image-only scanned PDFs still require OCR first.

Can I change Chapter 2 after I already wrote Chapter 5?

Yes. Edit Chapter 2 and click “Save & update memory”. Later generation will use the new Chapter 2 memory.

Do I need to regenerate the entire paper after changing Chapter 2?

No. Existing chapters are not rewritten automatically. Review only the chapters that may be affected by the change.

Will Lumielle create references that are not in my library?

Outline citations and chapter bindings are selected from literature records that already exist in the library, which reduces this risk. Final facts, quotations, and bibliography entries still require human checking.

Does a green model connection test mean the model will always work?

No. It confirms one successful call at that moment. Balance, network, provider status, and model availability can change later.

Can I submit the generated manuscript without review?

No. Lumielle is a writing assistant. The author remains responsible for research data, evidence, citations, ethics, academic rules, and the final text.

13

Troubleshooting

Problem	Try this first
Nothing opens after double-clicking the launcher	Wait a moment and keep the launcher window open. If the browser does not open automatically, check the launcher window for a local address.
macOS blocks the launcher	Control-click / right-click and choose Open; or allow the launch from System Settings → Privacy & Security.

Problem	Try this first
Fetching models fails	Check the Base URL, API Key, and network. DeepSeek users should confirm they are using an Open Platform API Key.
Model test says authentication failed	Paste the complete API Key again. If the key was lost, create a new one.
Model test says insufficient balance	Check the provider's Open Platform balance and top up if needed.
A literature PDF cannot be read	Check whether text can be selected in the PDF. Image-only scans need OCR first.
Literature AI analysis fails	The source remains imported and a previous valid profile is kept. Check the model connection and retry.
Word formatting does not match an institution	Check the selected preset. Special covers, headers, or bibliography rules may need an institution-specific preset or manual adjustment.

When asking for help

Keep a screenshot of the error and note the platform (Windows / macOS), the page you were using, and the steps that led to the issue. Do not expose your API Key in screenshots or support chats.

14

Data, privacy, and open source

Lumielle runs locally. Project settings, drafts, research material, literature records, and formatting presets are mainly stored in the data folder inside the installation directory. When an LLM task is run, the prompt and the relevant content needed for that task are sent to the model provider you selected.

Do not publicly share your working data folder

It may contain API configuration, drafts, literature, research material, and custom settings. If you need to share Lumielle with someone else, use a clean distribution package instead of copying your already-used working folder.

Lumielle is open source under the MIT License. Source code and issue reporting are available on GitHub.

GitHub: <https://github.com/Lumielle-BlueOMOcean/Lumielle-Academic-Assistant->

Community & feedback: QQ Group 1029688024.

15

Version history

v1.1.0 — Unreleased

Current development version. No final v1.1.0 tag or GitHub Release has been created yet.

Added

- An eight-section Research Fact Layer with complete-source parsing, editable smart filing suggestions, provenance, and chapter-writing controls.
- A planning digest that includes saved facts from every Research Foundation section; sections excluded from chapter writing can still influence planning.
- Complete-document literature analysis for PDF, DOCX, TXT, and pasted text, plus source-linked evidence restricted to literature bound to each chapter.
- Research-aware literature relevance; relevance ratings are separate from source-quality assessment.
- Versioned formatting presets with validated import/export, custom presets, model-generated previews, and deterministic Word rendering.
- Formatting preset support for grid / three-line tables, page-number starts, heading pagination controls, and figure-width limits; General Academic and Legacy Compatible use distinct layouts.
- A real minimal model-connection test that checks the API Key, selected model, and generation endpoint before use, with clear messages for authentication, balance, unavailable models, rate limits, and network errors.
- Rewritten Chinese and English user manuals for non-technical users while retaining the Lumielle poster.

Fixed

- Chapter Memory now processes the complete chapter instead of only the first 3,000 characters.
- A failed memory rebuild keeps the previous valid chapter memory.
- Clearing literature records and deleting imported source files are separate actions.
- Failed full-document literature analysis does not overwrite a previous valid profile or evidence store.
- Accepted smart-filing suggestions keep the editable role field.
- Continuous manual layout and the Lumielle poster were restored.

Changed

- Fresh DeepSeek profiles use <https://api.deepseek.com> and deepseek-flash; existing profiles are not migrated automatically, and the fetched model list remains authoritative.
- Context Sandwich documentation now describes complete-chapter memory and compact summaries without promising perfect recall or consistency.
- Logic planning uses a bounded digest covering all Research Foundation sections rather than a head-only prefix; chapter-writing toggles affect retrieval, not planning inclusion.
- Large literature libraries are reviewed in profile batches before hierarchical synthesis; relevance and source quality are assessed separately.
- Word export is driven by validated formatting presets while template-preserving mode keeps existing template content and styles.

v1.0.1 — 2026-09-27

This release focused on chapter context, citation stability, long-document handling, chart safety, LLM failure handling, manuscript reverse engineering, unreadable JSON protection, MIT licensing, Python 3.12 CI, and core regression tests.

- The Context Sandwich now constrains the actual chapter-generation request, and saved writing-style prompts are actually applied.
- Word export uses a global ordered reference registry and remaps chapter-local citations so the same number does not point to different sources.
- Charts use validated structured specifications and trusted local rendering instead of executing model-generated Python; illustrative data requires explicit permission.
- Research Foundation parsing and manuscript reverse engineering no longer use silent fixed-prefix truncation for long documents.
- Empty responses and API failures are blocked from drafts, outlines, and reports; unavailable literature rating is shown as Unrated.
- Unreadable/corrupted JSON files are preserved and protected from overwrite.
- MIT License, GitHub source, and QQ feedback information were aligned across the README, applications, and manuals.

v1.0.0 — 2026-08-26

Initial public release with Chinese and English desktop packages for Windows and macOS.